

Game

MinMax

Alpha-Beta

Motivation

- Why study games?
 - Games are fun!
 - Historical role in AI
 - Studying games teaches us how to deal with other agents trying to foil our plans
 - Huge state spaces – Games are hard!
 - Nice, clean environment with clear criteria for success

The Simple Case

- Chess, checkers, Tic-Tac-Toe, Othello, go ...
- Two players alternate moves
- **Zero-sum**: one player's loss is another's gain
- **Perfect Information**: each player knows the entire game state
- **Deterministic**: no element of chance
- Clear set of legal moves
- Well-defined outcomes (e.g. win, lose, draw)

More complicated games

- Most card games (e.g. Hearts, Bridge, “Belot” ...etc.) and Scrabble
 - non-deterministic
 - lacking in perfect information.
- Cooperative games

Game setup

- **Two players:** A and B
- **A** moves first and they take turns until the game is over. Winner gets award, loser gets penalty.
- **Games as search:**
 - **Initial state:** e.g. board configuration of chess
 - **Successor function:** list of (move,state) pairs specifying legal moves.
 - **Goal test:** Is the game finished?
 - **Utility function:** Gives numerical value of terminal states. E.g. win (+1), lose (-1) and draw (0) in tic-tac-toe
- **A** uses **search tree** to determine next move.

How to Play a Game by Searching

- **General Scheme**
 - Consider all legal moves, each of which will lead to some new state of the environment ('board position')
 - **Evaluate** each possible resulting board position
 - Pick the move which leads to the best board position.
 - Wait for your opponent's move, then **repeat**.
- **Key problems**
 - Representing the 'board'
 - Representing legal next boards
 - Evaluating positions
 - Looking ahead

Game Trees

- Represent the problem space for a game by a tree
 - **Nodes** represent 'board positions' (state)
 - **edges** represent legal moves.
- **Root node** is the position in which a decision must be made.
- **Evaluation function f** assigns real-number scores to 'board positions.'
- **Terminal nodes (leaf)** represent ways the game could end, labeled with the desirability of that ending (e.g. win/lose/draw or a numerical score)

MAX & MIN Nodes

- When I move, I attempt to **MAXimize** my performance.
- When my opponent moves, he attempts to **MINimize** my performance.

TO REPRESENT THIS:

- If we move first, label the root MAX; if our opponent does, label it MIN.
- Alternate labels for each successive tree level.
 - if the root (level 0) is our turn (MAX), all even levels will represent turns for us (MAX), and all odd ones turns for our opponent (MIN).

Evaluation functions

- Evaluations how good a 'board position' is
 - Based on **static features** of that board alone
- Zero-sum assumption lets us use one function to describe goodness for both players.
 - $f(n) > 0$ if we are winning in position n
 - $f(n) = 0$ if position n is tied
 - $f(n) < 0$ if our opponent is winning in position n
- Build using expert knowledge (**Heuristic**),
 - Tic-tac-toe: $f(n) = (\# \text{ of 3 lengths possible for me}) - (\# \text{ possible for you})$

Chess Evaluation Functions

- **Alan Turing's**
 $f(n) = (\text{sum of your piece values}) - (\text{sum of opponent's piece values})$

Pawn	1.0
Knight	3.0
Bishop	3.25
Rook	5.0
Queen	9.0

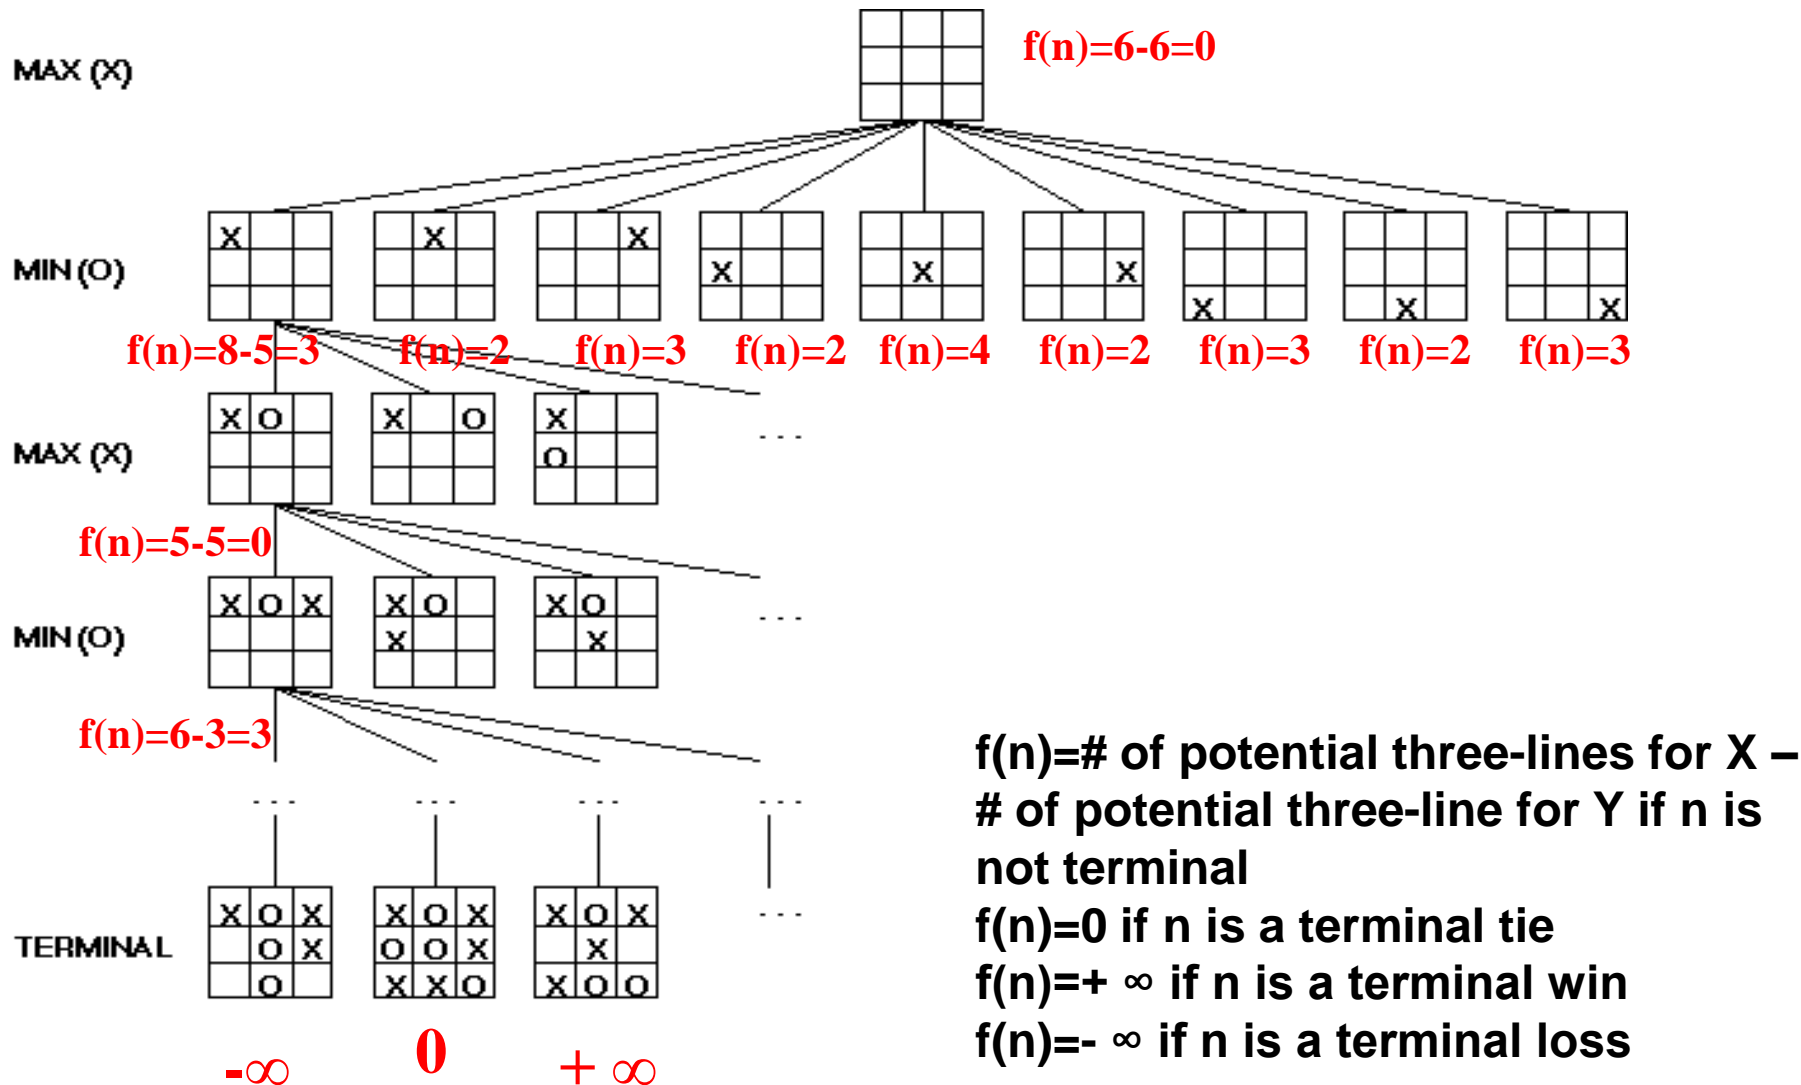
- More complex: weighted sum of positional features:

$$\sum w_i \text{feature}_i(n)$$


Pieces values for a simple
Turing-style evaluation
function

- Deep Blue has > 8000 features
(IBM Computer vs. Gary Kasparov)

A Partial Game Tree for Tic-Tac-Toe




Some Chess Positions and their Evaluations




White to move

$$f(n) = (9+3) - (5+5+3.25) \\ = -1.25$$



$$\dots Nxg5?? \\ f(n) = (9+3) - (5+5) \\ = 2$$



$$\text{Uh-oh: Rxg4+} \\ f(n) = (3) - (5+5) \\ = -7$$

So, considering our opponent's possible responses would be wise.


And black may force checkmate

MinMax Algorithm


- Two-player games with perfect information, the **minmax** can determine the best move for a player by enumerating (evaluating) the entire game tree.
- Player 1 is called Max
 - Maximizes result
- Player 2 is called Min
 - Minimizes opponent's result

MinMax Example


- Perfect play for deterministic games
-
- **Idea**: choose move to position with highest **minimax value**
= best achievable payoff against best play
-
- E.g., 2-ply game:



Two-Ply Game Tree



Two-Ply Game Tree



MinMax steps

```
Int MinMax (state s, int depth, int type)
{
    if( terminate(s)) return Eval(s);
    if( type == max )
    {
        for ( child =1, child <= NmbSuccessor(s); child++)
        {
            value = MinMax(Successor(s, child), depth+1, Min)
            if( value > BestScore) BestScore = Value;
        }
    }
    if( type == min )
    {
        for ( child =1, child <= NmbSuccessor(s); child++)
        {
            value = MinMax(Successor(s, child), depth+1, Max)
            if( value < BestScore) BestScore = Value;
        }
    }
    return BestScore;
}
```

MinMax Analysis


- Time Complexity: $O(b^d)$ ☹️
- Space Complexity: $O(b*d)$ 😊
- Optimality: Yes 😊

Problem: Game → Resources Limited!

– Time to make an action is limited

- Can we do better ? Yes !
- How ? Cutting useless branches !

Some nodes in the search can be *proven* to be irrelevant to the outcome of the search




How do we deal with resource limits?

- **Evaluation function**: return an estimate of the expected utility of the game from a given position, i.e.:
 - **Generate Search Tree up to certain level (i.e.: 4)**
- **Alpha-beta pruning**:
 - return appropriate minimax decision without exploring entire tree


Alpha-Beta Algorithm

- It is based on process of eliminating a branch of the search tree “pruning” the search tree.
- It is applied as standard minmax tree:
 - it returns the same move as minimax
 - prunes away branches that are not necessary to the final decision.

α - β Example



Alpha-Beta




MAX


$[-\infty, +\infty]$

MIN


$[-\infty, +\infty]$




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)



Alpha-Beta Example (continued)




Alpha-Beta Example (continued)




- We don't need to compute the value at this node.
- No matter what it is it can't effect the value of the root node.


Alpha-Beta Example (continued)




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)



Alpha-Beta Example (continued)



Example Alpha-Beta



Alpha-Beta Analysis

- In perfect case (perfect ordering) the depth is decreased twice in time complexity:
 - ➔ $O(b^{d/2})$
 - ➔ which means that the branching factor (b) is decreased to \sqrt{b}

Effectiveness of Alpha-Beta Pruning

- Guaranteed to compute same root value as Minimax
- **Worst case:** no pruning, same as Minimax ($O(b^d)$)
- **Best case:** when each player's best move is the first option examined, you examine only $O(b^{d/2})$ nodes, allowing you to search twice as deep!
- For **Deep Blue**, alpha-beta pruning reduced the average branching factor from 35-40 to 6.

Game Trees

- Represent the problem space for a game by a tree
 - **Nodes** represent 'board positions' (state)
 - **edges** represent legal moves.
- **Root node** is the position in which a decision must be made.
- **Evaluation function f** assigns real-number scores to 'board positions.'
- **Terminal nodes (leaf)** represent ways the game could end, labeled with the desirability of that ending (e.g. win/lose/draw or a numerical score)

MAX & MIN Nodes

- When I move, I attempt to **MAXimize** my performance.
- When my opponent moves, he attempts to **MINimize** my performance.

TO REPRESENT THIS:

- If we move first, label the root MAX; if our opponent does, label it MIN.
- Alternate labels for each successive tree level.
 - if the root (level 0) is our turn (MAX), all even levels will represent turns for us (MAX), and all odd ones turns for our opponent (MIN).

Evaluation functions

- Evaluations how good a 'board position' is
 - Based on **static features** of that board alone
- Zero-sum assumption lets us use one function to describe goodness for both players.
 - $f(n) > 0$ if we are winning in position n
 - $f(n) = 0$ if position n is tied
 - $f(n) < 0$ if our opponent is winning in position n
- Build using expert knowledge (**Heuristic**),
 - Tic-tac-toe: $f(n) = (\# \text{ of 3 lengths possible for me}) - (\# \text{ possible for you})$

Chess Evaluation Functions

- **Alan Turing's**
 $f(n) = (\text{sum of your piece values}) - (\text{sum of opponent's piece values})$

Pawn	1.0
Knight	3.0
Bishop	3.25
Rook	5.0
Queen	9.0

- More complex: weighted sum of positional features:

$$\sum w_i \text{feature}_i(n)$$

Pieces values for a simple
Turing-style evaluation
function

- Deep Blue has > 8000 features
(IBM Computer vs. Gary Kasparov)