

Diabetes Predictor using Logistic Regression

Consider below dataset

Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
6	148	72	35	0	33.6	0.627	50	1
1	85	66	29	0	26.6	0.351	31	0
8	183	64	0	0	23.3	0.672	32	1
1	89	66	23	94	28.1	0.167	21	0
0	137	40	35	168	43.1	2.288	33	1
5	116	74	0	0	25.6	0.201	30	0
3	78	50	32	88	31	0.248	26	1
10	115	0	0	0	35.3	0.134	29	0
2	197	70	45	543	30.5	0.158	53	1
8	125	96	0	0	0	0.232	54	1

Diabetes predictor application using Logistic Regression algorithm

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from warnings import simplefilter

simplefilter(action='ignore', category=FutureWarning)

print("---- Marvellous Infosystems by Piyush Khairnar----")

print("---- Diabetes predictor using Logistic Regression -----")

diabetes = pd.read_csv('diabetes.csv')

print("Columns of Dataset")
print(diabetes.columns)

print("First 5 records of dataset")
print(diabetes.head())

print("Dimension of diabetes data: {}".format(diabetes.shape))
```

```
X_train, X_test, y_train, y_test = train_test_split(diabetes.loc[:, diabetes.columns
!= 'Outcome'], diabetes['Outcome'], stratify=diabetes['Outcome'],
random_state=66)
```

```
logreg = LogisticRegression().fit(X_train, y_train)
```

```
print("Training set accuracy: {:.3f}".format(logreg.score(X_train, y_train)))
```

```
print("Test set accuracy: {:.3f}".format(logreg.score(X_test, y_test)))
```

```
logreg001 = LogisticRegression(C=0.01).fit(X_train, y_train)
```

```
print("Training set accuracy: {:.3f}".format(logreg001.score(X_train, y_train)))
```

```
print("Test set accuracy: {:.3f}".format(logreg001.score(X_test, y_test)))
```

Output of above application

```
---- Marvellous Infosystems by Piyush Khairnar----
```

```
---- Diabetes predictor using Logistic Regression -----
```

```
Columns of Dataset
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
      'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

```
First 5 records of dataset
```

	Pregnancies	Glucose	BloodPressure	...	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72 ...	0.627	50	1	
1	1	85	66 ...	0.351	31	0	
2	8	183	64 ...	0.672	32	1	
3	1	89	66 ...	0.167	21	0	
4	0	137	40 ...	2.288	33	1	

```
[5 rows x 9 columns]
```

```
Dimension of diabetes data: (768, 9)
```

```
Training set accuracy: 0.781
```

```
Test set accuracy: 0.771
```

```
Training set accuracy: 0.700
```

```
Test set accuracy: 0.703
```